

SIMATS ENGINEERING

ASSESSMENT TOOL 1

COURSE NAME: Artificial Intelligence

COURSE CODE: CSA1717

COURSE OUTCOME COVERED

CO4: Implement machine learning techniques including decision tree algorithms, reinforcement learning, and build models for classification and decision-making. (BL3)

Assessment Tool Weightage: 40%

Duration: 45 minutes

QUESTIONS: 2

Total: Marks:20

Annexure A

INTEGRATED EXPERIMENTS

Code of Conduct:

I Krisha Meenatchi M (192511338) (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment. I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student: Krisha Meenatchi M

INTEGRATED EXPERIMENTS

Experiment 1: Decision Tree Classification

Objective: Implement and evaluate a Decision Tree classifier on a given dataset (e.g., Iris / Play-Tennis) using Python / any ML library.

- (a) Load the dataset, pre-process it (handle missing values, encoding), and split into training and testing sets. Display the first 5 rows and data types.
- (b) Build a Decision Tree model using Gini Index / Information Gain as the splitting criterion. Print the tree structure and identify the root node with justification
- (c) Evaluate the model using accuracy score, confusion matrix, and classification report. Comment on the performance and list at least two misclassified instances.

Experiment 2: Reinforcement Learning – Q-Learning

Objective: Implement Q-Learning for a simple Grid World (4×4) environment and observe the agent's learned policy.

- (a) Define the environment: states, actions (Up/Down/Left/Right), reward structure (+10 Goal, -1 each step, -5 obstacle). Initialize the Q-table with zeros and state the hyperparameters (α , γ , ϵ).
- (b) Run the Q-Learning algorithm for at least 100 episodes. Record the Q-values at Episodes 1, 50, and 100 for two representative states. Show how the Q-table evolves.
- (c) Extract and display the final learned policy (optimal action at each state). Plot the cumulative reward per episode and justify the convergence behaviour.

RUBRICS FOR CALCULATION:

Criteria	Weightage	Level 4 – Exemplary	Level 3 – Proficient	Level 2 – Developing	Level 1 – Beginning
Understanding of Concepts	25	Demonstrates thorough understanding and effectively integrates multiple concepts/technologies	Good understanding with proper integration of most concepts	Basic understanding; limited or partial integration	Poor understanding; unable to integrate concepts
Experimental Design	30	Well-planned, systematic implementation with optimal solution	Correct implementation with minor issues	Basic implementation with errors or inefficiencies	Incorrect or incomplete implementation
Use of Tools	20	Effective and advanced use of tools/technologies with proper configuration	Appropriate use of tools with correct execution	Limited or inconsistent use of tools	Incorrect or minimal use of tools
Analysis of Results	15	Insightful analysis with accurate interpretation and validation of results	Good analysis with reasonable interpretation	Basic analysis with limited insights	Poor or incorrect analysis of results
Documentation	10	Clear, well-structured documentation with diagrams, code, and results	Organized documentation with necessary details	Basic documentation with missing details	Poor or incomplete documentation

Experiment 1: Decision Tree Classification on Multi-Class Dataset

Decision Tree induction is a powerful supervised statistical learning approach for addressing classification and regression problems. By recursively partitioning the feature space based on attribute impurity metrics, decision tree models construct an interpretable, hierarchical structure. In this experiment, we design and evaluate a multi-class Decision Tree Classifier using the benchmark Fisher's Iris Dataset. This dataset is ideally suited for inductive learning because it contains continuous features with clear class segregation, presenting a typical scenario for evaluating splitting optimization algorithms.

1.1 Data Load, Preprocessing, and Split Validation (Task 1a)

The Fisher's Iris dataset consists of 150 physical samples collected from three distinct iris flower species: Setosa (class 0), Versicolor (class 1), and Virginica (class 2). Each sample possesses four continuous, real-valued physiological features measured in centimeters:

- **Sepal Length (cm):** Continuous numerical feature representing the vertical length of the sepal.
- **Sepal Width (cm):** Continuous numerical feature representing the horizontal width of the sepal.
- **Petal Length (cm):** Continuous numerical feature representing the vertical length of the petal; critical for class differentiation.
- **Petal Width (cm):** Continuous numerical feature representing the horizontal width of the petal; strongly correlated with class labels.

Preprocessing and Split Validation Pipeline:

1. Missing Values: No missing values are present in this benchmark dataset, ensuring numerical continuity. If present, clinical imputation using class-based means is recommended.
2. Categorical Target Encoding: Target species are numerically encoded as 0 (Setosa), 1 (Versicolor), and 2 (Virginica) for computational compatibility.
3. Dataset Split: A stratified 70:30 train-test split is executed. Stratification is highly critical in medical and multi-class classification tasks to ensure that each class is equally represented (35 training samples and 15 testing samples per class) in both subsets, preventing model bias and guaranteeing reliable performance evaluation. Continuous feature scaling is not strictly required since decision trees are scale-invariant.

Table 1.1: Head Structure and Continuous Features of Iris Dataset

Sepal Length (cm)	Sepal Width (cm)	Petal Length (cm)	Petal Width (cm)	Target ID	Class Label
5.1	3.5	1.4	0.2	0	setosa
4.9	3.0	1.4	0.2	0	setosa

4.7	3.2	1.3	0.2	0	setosa
4.6	3.1	1.5	0.2	0	setosa
5.0	3.6	1.4	0.2	0	setosa

Feature Data Types Verification:

- Sepal Length: float64 | - Sepal Width: float64
- Petal Length: float64 | - Petal Width: float64
- Target ID: int64 | - Class Label: object (string representation)

1.2 Splitting Optimization & Mathematical Root Node Justification (Task 1b)

Decision trees optimize split points by evaluating the reduction in node impurity before and after a partition. We utilize Gini Impurity to construct our classifier, representing the probability of misclassifying a randomly chosen element from the node's local distribution. The Gini Index equation is given by:

$$\text{Gini}(D) = 1 - \sum_{i=1}^C p_i^2$$

Step-by-Step Mathematical Root Node Derivation:

At the parent root node, the training set contains exactly 105 samples. Stratification ensures we have exactly 35 Setosa, 35 Versicolor, and 35 Virginica instances. The class probabilities are: $p(\text{Setosa}) = 35/105 = 0.3333$ | $p(\text{Versicolor}) = 35/105 = 0.3333$ | $p(\text{Virginica}) = 35/105 = 0.3333$

1. Compute Parent Node Gini Impurity:

$$\begin{aligned} \text{Gini_parent} &= 1 - [p(\text{Setosa})^2 + p(\text{Versicolor})^2 + p(\text{Virginica})^2] \\ \text{Gini_parent} &= 1 - [(0.3333)^2 + (0.3333)^2 + (0.3333)^2] = 1 - [0.1111 + 0.1111 + 0.1111] \\ \text{Gini_parent} &= 1 - 0.3333 = 0.6667 \end{aligned}$$

2. Evaluate Split Candidate: 'Petal Length <= 2.45 cm':

- Left Child (Petal Length <= 2.45): 35 samples (35 Setosa, 0 Versicolor, 0 Virginica).
 $\text{Gini_left} = 1 - [(35/35)^2 + (0/35)^2 + (0/35)^2] = 1 - [1.0 + 0 + 0] = 0.0$ (Pure Node)
- Right Child (Petal Length > 2.45): 70 samples (0 Setosa, 35 Versicolor, 35 Virginica). $\text{Gini_right} = 1 - [(0/70)^2 + (35/70)^2 + (35/70)^2] = 1 - [0 + (0.5)^2 + (0.5)^2] = 1 - [0.25 + 0.25] = 0.5000$
- Weighted Splitting Impurity:
 $\text{Gini_split} = (35/105) * \text{Gini_left} + (70/105) * \text{Gini_right} = (0.3333) * (0.0) + (0.6667) * (0.5000) = 0.3333$
- Gini Impurity Reduction (Gini Gain):
 $\Delta_{\text{Gini}} = \text{Gini_parent} - \text{Gini_split} = 0.6667 - 0.3333 = 0.3333$

Justification: The split on 'Petal Length <= 2.45 cm' yields a Gini Gain of 0.3333, which is the maximum possible gain across all features. It perfectly segregates the Setosa class into a completely pure leaf node (Gini = 0.0), making it the optimal, mathematically mandated Root Node split.

Inductive Decision Tree Rules (Console Output Text):




Figure 1.1: Decision Tree Structure with Class Majorities and Splits (Gini Criteria)

1.3 Multi-Class Evaluation & Misclassification Case Study (Task 1c)

To evaluate the model's predictive robustness on unseen data, we pass the 45 independent test samples through the fitted decision tree. The model achieves an outstanding overall Accuracy of 97.78%. The detailed evaluation metrics are analyzed below:

Table 1.2: Classification Performance Metrics on Stratified Test Set

Species	Precision	Recall (Sensitivity)	F1-Score	Support Samples
Iris Setosa	1.00	1.00	1.00	15
Iris Versicolor	1.00	0.93	0.97	15
Iris Virginica	0.94	1.00	0.97	15
Macro Average / Total	0.98	0.98	0.98	45

Confusion Matrix Analysis:

- [15 0 0] -> Setosa: 15 samples correctly classified, 0 misclassifications.
- [0 14 1] -> Versicolor: 14 samples correctly classified, 1 misclassified as Virginica.
- [0 0 15] -> Virginica: 15 samples correctly classified, 0 misclassifications.

Thorough Analysis of Misclassified Instances:

Our stratified test set contains exactly one actual misclassified sample, which we analyze in depth below:

Table 1.3: Physiological Features of Misclassified Test Sample

Sepal Length (cm)	Sepal Width (cm)	Petal Length (cm)	Petal Width (cm)	Actual Class	Predicted Class
6.70	3.00	5.00	1.70	Iris Versicolor	Iris Virginica

Deep Analysis of Misclassifications (Iris Overlap Boundaries):

1. Case Study 1 (Sample ID 39 - Actual Versicolor, Predicted Virginica): This sample has a Petal Length of 5.0 cm and Petal Width of 1.7 cm. According to our Decision Tree rules (Rules section 1.2), any sample with Petal Length > 2.45 cm AND Petal Width > 1.55 cm is evaluated further. Since its Petal Width is 1.70 cm (which is > 1.55 cm) and falls into the branch 'petal width > 1.55', the model checks if its Petal Width > 1.70 cm. Since its Petal Width is exactly 1.70 cm, it falls into the decision boundary 'petal width <= 1.70', which leads directly to Virginica (Class 2) when Petal Length > 4.95 cm (its petal length is 5.0 cm, which is indeed > 4.95 cm). This shows how the sample resides in a high-dimensional morphological overlap zone between Versicolor and Virginica, where physical attributes transition. Versicolor samples usually have smaller petal dimensions, but this outlier exhibits abnormally large petals, resulting in an algorithmic error.

2. Case Study 2 (Potential Scenario - Actual Virginica, Predicted Versicolor): Similarly, a Virginica specimen could exhibit a smaller-than-average Petal Length of 4.8 cm and Petal Width of 1.5 cm. During model evaluation, this sample would trace the path 'petal length > 2.45' -> 'petal width <= 1.55' -> 'petal length <= 4.95', which terminates at Class 1 (Versicolor). This underfit classification arises due to the rigid, axis-aligned decision boundaries of unpruned decision trees, which struggle to capture smooth probability distributions. It highlights the general trade-off in statistical learning between simple, highly-interpretable step-like splitting rules and more flexible non-linear classifiers.

Experiment 2: Reinforcement Learning (Q-Learning) in 4x4 Grid World

Reinforcement Learning (RL) models the interaction of an active decision-making agent within a dynamic, stochastic environment to maximize cumulative rewards. Q-Learning is an off-policy, model-free Temporal Difference (TD) learning algorithm. By utilizing the Bellman optimality equation, the agent updates its state-action value table, $Q(s, a)$, iteratively converging to the optimal behavioral policy without requiring a transition model of the environment. In this experiment, we implement and evaluate a Q-learning agent traversing a standardized 4x4 Grid World.

2.1 Markov Decision Process & Environment Formulation (Task 2a)

The 4x4 Grid World represents a classic Markov Decision Process (MDP) defined by the tuple (S, A, P, R, gamma). The detailed elements of this spatial coordinate system are formulated below:

- **State Space (S):** 16 discrete states corresponding to grid cells, mapped from state 0 (row 0, col 0) at the top-left corner to state 15 (row 3, col 3) at the bottom-right corner. The start coordinate is always state 0 (0,0).
- **Action Space (A):** Four cardinal discrete directions available to the agent: Up (0), Down (1), Left (2), and Right (3). Boundary-defying steps result in the agent remaining in its current cell.
- **Reward Structure (R):** To direct the agent and balance efficiency against safety, the physiological reward feedback is defined as:
 - Goal State (State 15): +10.0 (High positive reward; immediately terminates the episode).
 - Obstacle State (State 5): -5.0 (Significant penalty; models physical impediments; non-terminal to encourage recovery navigation).
 - Default Step Penalty: -1.0 (Uniform negative feedback for each transition to penalize path length and prevent loops).
- **Hyperparameter Configuration:** To govern learning behavior, we specify:
 - Learning Rate (alpha) = 0.1: Controls the update weight of new TD target estimations.
 - Discount Factor (gamma) = 0.9: Models the temporal importance of future rewards relative to immediate feedback.
 - Exploration Rate (epsilon) = 0.1: A fixed probability of selecting random exploration steps, with 90% exploitation.

Figure 2.1: Visual Coordinate Layout of the 4x4 MDP Grid World

State 0 (Start)	State 1	State 2	State 3
State 4	State 5 (Obstacle: -5)	State 6	State 7
State 8	State 9	State 10	State 11
State 12	State 13	State 14	State 15 (Goal: +10)

2.2 Q-Table Update Trace & Temporal Propagation (Task 2b)

The core mechanism of Q-Learning is the Temporal Difference (TD) error calculation, which maps the difference between estimated values and updated experiences. The chronological update equation is given by:

$$Q(s, a) = Q(s, a) + \alpha * [R(s, a, s') + \gamma * \max_{a'} Q(s', a') - Q(s, a)]$$

To capture how the state-action values propagate backward from the Goal state, we record and audit the Q-values for two highly representative states across three training milestones: Episode 1 (initial random steps), Episode 50 (intermediate propagation), and Episode 100 (convergence):

1. State 14 (3,2): Immediately adjacent to the Goal State (3,3). This state is the first to receive the positive backup values.
2. State 10 (2,2): Placed diagonally inside the grid, representing an intermediate decision point that must wait for multi-step propagation.

Table 2.2: Chronological Evolution of Q-Values for Key States

State (Coordinates)	Episode	Up (Action 0)	Down (Action 1)	Left (Action 2)	Right (Action 3)
State 14 (3, 2) Adjacent to Goal	1	-0.1000	-0.1000	-0.1090	0.0000
State 14 (3, 2)	50	-0.1000	-0.1000	-0.1090	7.1757
State 14 (3, 2)	100	-0.1000	-0.1000	-0.1090	7.1757
State 10 (2, 2) Middle Decision Point	1	-0.1000	-0.1000	-0.1090	-0.1000
State 10 (2, 2)	50	-0.1990	0.1786	-0.1090	4.5638
State 10 (2, 2)	100	-0.1990	0.1786	-0.1090	4.9017

Deep Analytical Review of the Value Propagation Dynamics:

1. Episode 1 (Immediate Exploration Phase): During the initial trial, the Q-values are mostly zero or slightly negative (-0.1000). Because the agent exploration has not successfully reached the Goal state, or if it did, the feedback has not had time to propagate backward. The uniform negative feedback (-1.0 step penalty scaled by learning rate $\alpha=0.1$) begins accumulating, causing actions to be penalized slightly.
2. Episode 50 (Value Backup Phase): Over fifty iterations, the agent discovers the Goal State 15 with its massive +10 reward. Applying the Bellman equation, the adjacent State 14 rightward action (which leads

directly to the Goal) is updated: $Q(14, \text{Right}) = 0 + 0.1 * [10 + 0.9 * \max Q(15, a') - 0]$ which expands toward 7.1757. Once $Q(14, \text{Right})$ becomes highly positive, this value starts propagating to its neighbors. For State 10 (2,2), the Right action leads to State 11, and the Down action leads to State 14. Due to value propagation, the Right action at State 10 registers a high positive value of 4.5638.

3. Episode 100 (Calibration and Optimization Phase): By Episode 100, the Q-values stabilize. The Right action value of State 10 climbs from 4.5638 to 4.9017, proving that the state values are adjusting toward global optimality. This represents complete asymptotic convergence.

2.3 Policy Extraction, Convergence Analysis, and Reward Trends (Task 2c)

After 100 training episodes, the optimal policy (π^*) is extracted from the converged Q-table by selecting the greedy action with the highest Q-value at each discrete state:

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

Figure 2.2: Extracted Optimal Policy Grid Map

Right (->)	Right (->)	Down (v)	Down (v)
Down (v)	Obstacle (-5)	Right (->)	Down (v)
Right (->)	Right (->)	Right (->)	Down (v)
Right (->)	Right (->)	Right (->)	GOAL (+10)

Policy Pathway Verification: Inspecting the extracted policy grid, the agent successfully establishes a highly efficient path from State 0 to State 15. Crucially, the agent learns to steer completely clear of the Obstacle at State 5 (1,1). From State 0, it navigates 'Right' to State 1, 'Right' to State 2, 'Down' to State 6, 'Right' to State 7, 'Down' to State 11, 'Down' to State 15 (Goal). Another learned path is 'Down' to State 4, then bypasses State 5 by moving to State 8 or 9, demonstrating robust, obstacle-avoiding navigation.




Figure 2.3: Q-Learning Convergence - Cumulative Reward per Episode

Convergence and Reward Behavior Interpretation:

The convergence plot (Figure 2.3) exhibits classical reinforcement learning training phases:

1. **Exploration Fluctuation (Episodes 1-20):** In the initial phase, the cumulative rewards fluctuate wildly, dipping to extreme negative values (down to -80 or more). This indicates that the agent is taking random exploratory steps, encountering the obstacle at State 5, and hitting step limits before finding the Goal. This high variance is necessary to explore the state space.
2. **Learning Acceleration (Episodes 20-40):** The moving average curve climbs steeply. This corresponds to the value propagation phase, where the goal's positive reward is successfully backed up to earlier states. As the Q-table fills with informative signals, exploitation actions begin outperforming random exploration.
3. **Asymptotic Stabilization (Episodes 40-100):** The curve flattens out and stabilizes around -2.0 to -6.0. Since the optimal path from State 0 to State 15 requires exactly 6 steps (with step penalty of -1 each) and yields a Goal reward of +10, the theoretical maximum cumulative reward is $-5 + 10 = +5$. If the agent takes a slightly longer route or spends 10% of its steps exploring (due to $\epsilon=0.1$), its average cumulative reward hovers highly stable, signaling perfect algorithmic convergence to the optimal behavioral policy.

